

ARCHIT AGRAWAL

Tempe, AZ 623-312-0435 architagrawal000@gmail.com linkedin.com/in/agrawal-archit github.com/architagrawal agrawal-archit.vercel.app

EXPERIENCE

AI Software Engineer, EdPlus, Arizona State University

Tempe, AZ May 2026 – Present

- Architected the agentic survey-analysis platform that retired every per-survey pipeline: 10 agents over 65 tools in TypeScript/NestJS on a Step Functions machine generated from a declarative graph, with Distributed Map fan-out and a 5% failure circuit breaker.
- Eliminated a silent failure mode dropping roughly 1 row in 8 unflagged: schema-constrained decoding, correlation tokens and citation checks took join integrity to 100% and label churn to zero across 258 runs.
- Isolated the dominant variable in labeling quality: model choice moved F1 by 0.230 against 0.007 for pipeline shape, 33x, redirecting the team's roadmap from orchestration to model selection.
- Retired the 0.637 accuracy baseline every prior claim rested on, after fresh runs reproduced it at 0.523 and traced the gap to a spreadsheet human-reviewed before delivery.
- Ran the platform with no database, on in-process DuckDB over per-run JSON with Lance retrieval from S3: ask latency held flat at 200x row volume, cross-scope queries over 216 runs in milliseconds.
- Authored the target architecture the production platform is built from: 80 logged decisions and 40 measured experiments, each stating what it does not establish.

Founding AI/ML Engineer, MyStage Music Inc.

Remote Jul 2025 – May 2026

- Owned the backend as founding engineer: re-architected a 14-Cloud-Function pipeline into one checkpointed LangGraph agent-worker on Cloud Run, cutting production service dependencies 64%, from 14 to 5.
- Built and operate a Playwright ingestion service indexing 70,000+ records a day across 650+ venues, with proxy rotation and atomic Firestore claim transactions making multi-worker retries idempotent.
- Lifted downstream search accuracy 25% with Gemini entity resolution and dedupe applied before records reach Firestore and Algolia.
- Delivered sub-50 ms search behind FastAPI and Algolia with real-time Firestore sync, traced by Logfire propagation that survives pause and resume, covered by 200+ pytest tests.

Software Engineer, EdPlus, Arizona State University

Tempe, AZ Sep 2023 – May 2025

- Led end-to-end development of a multi-tenant RAG assistant used by 1,000+ faculty to author courses reaching 60,000+ students.
- Made transcript analysis 16x faster, 4 hours to 15 minutes, pairing a Neo4j knowledge graph with LLM-generated Cypher constrained to a schema allow-list.
- Enforced tenant isolation at retrieval rather than in the prompt, so one college's material cannot surface in another's answer, and made onboarding a college a config change, not a deployment.
- Shipped prompt changes on evidence with a Prompt Flow harness scoring groundedness, relevance and coherence on a fixed question set, plus hybrid retrieval for the course codes an embedding blurs.

Student Researcher, Arizona State University

Tempe, AZ Aug 2024 – May 2025

- Built a PyTorch and Diffusers harness for text-to-video generation scoring physical plausibility separately from semantic fidelity, holding seed, resolution and sampler fixed so a score gap measures the model.
- Showed fluency and physical plausibility correlate only weakly, so a model ranked first on human preference can rank last on causality.

Data Research Aide, Knowledge Exchange for Resilience, Arizona State University

Tempe, AZ Jun 2024 – Aug 2024

- Cut p95 retrieval latency 60% at 95% recall with a hybrid FAISS and Neo4j layer, choosing the index by measurement and stating the recall cost of leaving an exact one.
- Surfaced collaborator recommendations at 85% top-k precision from a FastAPI embedding service over Postgres/pgvector, fed by an idempotent ETL ingesting 10,000+ profiles a day, plus 20+ REST APIs on .NET 8 with Dapper, MediatR and SQS.

Software Engineer, Zeus Learning

Mumbai, India Jan 2022 – Jul 2023

- Shipped a Redis-backed desk-demand prediction service across 300+ Fortune 500 sites including Goldman Sachs and Merck, improving reported occupancy 30%.
- Split a .NET monolith into Kubernetes microservices as a strangler migration behind the existing API, no big-bang cutover; footprint down 35%, infrastructure cost down 20%.
- Cut deploy time 70% and production incidents 40% by making the SonarQube gate block the merge, scoped to new code so a legacy backlog could not neuter it on day one.

Product Intern, EAT.FIT

Bengaluru, India Sep 2021 – Dec 2021

- Cut driver-location serving cost 45% by keying poll intervals to distance remaining, interpolating client-side so slower polling still rendered as continuous movement.

EDUCATION

M.S., Computer Science, Arizona State University, GPA 4.0

Aug 2023 – May 2025

B.Tech., Information and Communication Technology, DA-IICT

Aug 2018 – May 2022